

Technical Documentation
NewsBot Intelligence System 2.0

Jose Moreno
ITAI 2373 - Natural Language Processing
Final Project

1. Introduction

NewsBot Intelligence System 2.0 is an end-to-end Natural Language Processing platform developed as the final project for ITAI 2373. The system extends the original midterm NewsBot instead of replacing it. The completed midterm foundation provides text preprocessing, TF-IDF feature extraction, linguistic analysis, sentiment analysis, supervised news classification, and Named Entity Recognition. The final-project extensions add advanced topic modeling, summarization, semantic search, multilingual processing, cross-language retrieval, and a conversational interface.

The purpose of the project is to demonstrate how classical NLP, machine learning, and newer language-model techniques can be integrated into one workflow for analyzing news articles. The system is implemented in a Jupyter/Google Colab notebook and uses the BBC News Train.csv dataset as its primary source of labeled news articles.

2. System Architecture

The NewsBot architecture follows a sequential but reusable NLP pipeline. Data is loaded and validated first, then cleaned and transformed into machine-readable features. The resulting representations support linguistic analysis and machine-learning classification. Advanced modules reuse the same article collection for topic discovery, summarization, semantic retrieval, multilingual analysis, and conversational exploration.

- Data layer: BBC News Train.csv is loaded into a Pandas DataFrame and inspected for shape, columns, category distribution, and missing values.
- Preprocessing layer: article text is normalized and prepared for downstream analysis.
- Feature layer: TF-IDF and count-based representations convert text into numerical vectors.
- Linguistic layer: POS tagging, dependency parsing, sentiment analysis, and Named Entity Recognition provide grammatical, emotional, and entity-level information.
- Machine-learning layer: multiple classifiers are trained and compared using held-out accuracy and cross-validation.
- Advanced analysis layer: LDA and NMF discover latent topics; summarization condenses articles; semantic search retrieves related content.

- Multilingual layer: language detection, translation, and cross-language search provide multilingual access.
- Interaction layer: a conversational query processor exposes statistics, search, sentiment, entities, topics, summaries, and translation through natural-language requests.
- Integration layer: NewsBotIntelligenceSystem2 combines the major capabilities into an end-to-end article analysis workflow.

3. Dataset

The project uses BBC News Train.csv. The dataset contains labeled BBC news articles from multiple categories and is used for exploratory analysis, model training, evaluation, topic modeling, search demonstrations, and integrated system tests. The notebook explicitly checks dataset dimensions, column names, category distribution, and missing values before processing.

A limitation of the dataset is that it represents one publisher and a limited historical collection. The final notebook also documents that authentic publication dates are not available for reliable temporal trend analysis. For that reason, the project does not claim genuine topic evolution over time.

4. Installation Guide

The project is designed to run in Google Colab. The notebook contains installation cells for the principal NLP and visualization packages and downloads the English spaCy model and required NLTK resources.

Recommended setup procedure:

- Open the completed NewsBot Intelligence System 2.0 notebook in Google Colab.
- Run the notebook from the beginning.
- When prompted by the upload cell, upload BBC News Train.csv.
- Allow the package-installation cell to install spaCy, scikit-learn, NLTK, Pandas, Matplotlib, Seaborn, WordCloud, and Plotly.
- Allow the notebook to download en_core_web_sm and the required NLTK corpora/resources.

- Continue running the remaining cells in order so that trained models and shared variables are available to later modules.
- Verify that the final evaluation table and completion checklist are executed without errors.

5. Core NLP Pipeline

5.1 Text Preprocessing

The preprocessing stage prepares raw article text for machine learning and linguistic analysis. Normalization reduces noise and creates a consistent representation for later modules. The preprocessing work from the midterm remains part of the final system so that the final extensions build an already tested foundation.

5.2 TF-IDF and Feature Engineering

TF-IDF converts article text into numerical feature vectors by weighting terms according to their importance within individual documents and across the corpus. These features support supervised classification and exploratory analysis. The project also uses count-based representations where appropriate for topic modeling.

5.3 POS Tagging and Dependency Parsing

NLTK and spaCy are used for grammatical analysis. Part-of-Speech tagging identifies grammatical roles, while dependency parsing represents relationships between words. These modules provide structural information that complements statistical text representations.

5.4 Sentiment Analysis

The system uses NLTK VADER sentiment analysis to estimate emotional tone. Sentiment features can be examined independently and incorporated into the broader analysis pipeline. The final documentation recognizes that lexicon-based sentiment methods may misread irony, complicated political language, or objectively written reports describing negative events.

5.5 Named Entity Recognition

spaCy Named Entity Recognition extracts entities such as people, organizations, locations, dates, and monetary references. Entity extraction supports richer article

analysis and business-intelligence use cases such as media monitoring and entity tracking.

6. Supervised News Classification

The classification system compares Multinomial Naive Bayes, Logistic Regression, and Support Vector Machine classifiers. Models are trained using the prepared feature matrix and evaluated against a held-out test set. Three-fold cross-validation is also used to provide a more stable estimate of model performance. The notebook automatically identifies the best-performing classifier according to test accuracy and reuses the trained model in the integrated system.

The final evaluation table reports the held-out classification accuracy produced by the executed notebook. Because the exact value is generated during execution, this technical report does not invent or replace that result; the executed notebook and its PDF preserve the actual run output.

7. Advanced Topic Modeling

The final-project extension implements both Latent Dirichlet Allocation (LDA) and Non-negative Matrix Factorization (NMF). These unsupervised methods discover latent themes from article text without relying directly on the known BBC category labels. The implementation includes topic-word inspection, topic visualization, article-topic assignment, and comparison between discovered topics and known categories.

Normalized Mutual Information is used in the final evaluation to compare topic assignments with known article categories. Both LDA and NMF results are reported by the executed notebook. Topic labels still require human interpretation because unsupervised topic models return distributions and high-weight terms rather than definitive human-readable labels.

8. Language Understanding and Generation

8.1 Summarization

NewsBot 2.0 supports article summarization using an extractive method and a transformer-based option with fallback behavior. Extractive summarization selects

important material from the source article, while the transformer path can produce a more natural condensed representation when the model is available. The fallback design helps keep the notebook usable when advanced model resources are unavailable.

8.2 Semantic Search

Semantic search allows the system to retrieve articles according to similarity of meaning rather than exact keyword overlap. The final evaluation records which semantic backend was successfully used during execution. This feature supports exploratory information retrieval and is combined with query expansion to improve the range of relevant results.

8.3 Query Expansion

Query expansion supplements the user's original request with related terms. The purpose is to reduce the limitations of literal keyword matching and improve retrieval when relevant articles use different wording.

9. Multilingual Intelligence

The multilingual module implements automatic language detection, translation, and cross-language search. The completed workflow demonstrates English-Spanish translation and allows translated or cross-language content to participate in retrieval and analysis. This module is intended to demonstrate multilingual access rather than claim perfect translation quality.

The notebook explicitly documents multilingual limitations. Automated translation may lose idioms, tone, regional meaning, or cultural context. Cross-language similarity can also vary depending on the available model and translation quality.

10. Conversational Interface

The conversational component provides a natural-language interface over the NewsBot analysis functions. The final evaluation identifies supported intents for statistics, search, sentiment, entities, topics, summaries, and translation. The interface classifies the user's intent and routes the request to the appropriate system capability.

Basic conversation-state management allows the interface to behave as an integrated exploration tool rather than a set of disconnected notebook functions. The design demonstrates how multiple NLP components can be exposed through a simpler user-facing interaction layer.

11. Integrated NewsBotIntelligenceSystem2

The final system of integration combines classification, confidence scoring, sentiment, entity analysis, topic identification, summarization, semantic search, language detection, translation, and conversational queries. In the end-to-end demonstration, test articles are analyzed and the system prints the predicted category, confidence, sentiment, topics, summary, and Spanish translation of the summary.

This integrated demonstration is important because it verifies that the final-project modules do not exist only as isolated experiments. They can operate together as part of a single NewsBot workflow.

12. API and Component Reference

Component	Purpose	Inputs / Outputs
Classifier	Predicts BBC news category and confidence.	Article features/text -> category, confidence
Sentiment analyzer	Estimates emotional tone with VADER.	Text -> sentiment label and compound score
NER pipeline	Extracts named entities with spaCy.	Text -> entity text/type pairs
LDA topic model	Discovers probabilistic latent topics.	Document-term representation -> topic distributions/words
NMF topic model	Discovers additive latent themes.	TF-IDF representation -> topic assignments/words
Summarizer	Produces concise article summaries.	Article text -> extractive or transformer/fallback summary

Semantic search	Finds articles related by meaning.	Natural-language query -> ranked articles
Language detector	Identifies text language.	Text -> detected language
Translation workflow	Translates content for multilingual access.	Text + target language -> translated text
Conversational processor	Routes natural language requests to system functions.	User query -> generated response
NewsBotIntelligenceSystem2	Coordinates the major final-project capabilities.	Title/content/target language -> integrated analysis result

13. Configuration Manual

The notebook is designed to be customized through parameters in its model and processing cells. The most important configuration choices include the number of TF-IDF features, train/test split settings, classifier hyperparameters, number of LDA/NMF topics, number of topic words displayed, summarization length or backend, semantic-search result count, translation target language, and conversational intent rules.

For reproducibility, `random_state=42` is used where applicable in machine-learning components. Users modifying the notebook should change one configuration at a time and rerun dependent cells so that trained models and evaluation results remain synchronized.

14. Testing and Evaluation

The completed notebook includes component-level and end-to-end evaluation. Classification uses held-out accuracy and cross-validation. LDA and NMF are evaluated with Normalized Mutual Information. Semantic search reports its active backend. Summarization reports on transformer/fallback status. Multilingual processing is demonstrated through detection, translation, and cross-language search. The conversational interface is checked against its supported intents.

The final completion checklist marks the following requirements complete: Modules 1-8 foundation, LDA, NMF, topic visualization, summarization, semantic search, query expansion, language detection, translation, cross-language search, conversational interface, NewsBot 2.0 integration, evaluation, and ethics/business-value discussion.

15. Limitations and Ethical Considerations

- The BBC dataset represents one publisher and a limited historical period.
- VADER can misinterpret irony, complex political language, and neutral reporting about negative events.
- Named Entity Recognition can miss uncommon names or misclassify ambiguous words.
- Automated translation may lose idioms, tone, and regional context.
- Extractive summaries can omit important context, while transformer summaries may introduce unsupported wording.
- The dataset does not provide authentic dates suitable for genuine temporal trend claims.
- Topic labels require human interpretation.
- Human review remains necessary for high-impact editorial, business, or public-facing decisions.

16. Business Value

NewsBot 2.0 demonstrates practical value for media monitoring, content routing, market research, multilingual access, semantic search, executive summarization, entity tracking, topic discovery, and interactive exploration. A media organization could use similar components to organize incoming stories automatically. Analysts could use entity and sentiment information to monitor coverage, while semantic retrieval and summarization could reduce the time required to locate and review relevant material.

17. Future Improvements

- Connect the system to live news feeds for real-time analysis.
- Use larger or domain-specific multilingual datasets.
- Fine-tuned transformer models for news classification and summarization.

- Improve conversational memory and support more complex follow-up questions.
- Add richer visual dashboards and exportable reports.
- Package notebook functions into reusable Python modules or a REST API for deployment.
- Add systematic bias detection, explainability, and more extensive automated testing.

18. Conclusion

NewsBot Intelligence System 2.0 demonstrates the integration of classical NLP, supervised machine learning, unsupervised topic modeling, semantic retrieval, language-model summarization, multilingual processing, and conversational AI in a single educational system. The final project preserves the working midterm foundation while extending it with the advanced techniques required for the course. The completed notebook and executed outputs demonstrate that the components can operate together in an end-to-end workflow and provide a strong foundation for future deployment-oriented development.